

---

000  
001  
002  
003  
004  
005  
006  
007 **BSMRL: A RISK-SENSITIVE AND TIME-AWARE**  
008 **FRAMEWORK FOR REINFORCEMENT LEARNING**  
009 **(DRAFT VER.)**

010  
011 **Anonymous authors**  
012 Paper under double-blind review

013 **ABSTRACT**  
014

015 In this paper, we propose BSMRL, a novel reinforcement learning frame-  
016 work that leverages the Black-Scholes-Merton (BSM) model for reward  
017 shaping. By integrating financial mathematics into the RL paradigm,  
018 BSMRL aims to enhance learning efficiency and stability in environments  
019 with sparse and delayed rewards. We further extend our approach by incor-  
020 porating the Merton Jump-Diffusion Model to account for sudden changes  
021 in the environment, providing a risk-sensitive and time-aware framework.  
022 Experimental results demonstrate that BSMRL outperforms traditional RL  
023 methods in various benchmark tasks, showcasing its potential for broader  
024 applications in software testing, quantitative finance, embodied AI, and  
025 beyond.

026 **CONTENTS**  
027  
028  
029  
030  
031  
032  
033  
034  
035  
036  
037  
038  
039  
040  
041  
042  
043  
044  
045  
046  
047  
048  
049  
050  
051  
052  
053

---

054    1 INTRODUCTION  
055

056    In the original Temporal-Difference (TD) learning framework, the agent learns to estimate  
057    the value function based on one-step transitions. However, in real-world scenarios, rewards  
058    can be sparse and delayed, making it challenging for agents to learn effective policies.  
059

060    Here, we propose BSMRL, a novel reinforcement learning framework that leverages the  
061    Black-Scholes-Merton (BSM) model for reward shaping. By integrating financial mathe-  
062    matics into the RL paradigm, BSMRL aims to enhance learning efficiency and stability in  
063    environments with sparse and delayed rewards. The BSM model provides a closed-form  
064    solution for European-style options, which can be adapted to shape rewards based on the  
065    agent's current state and time-to-go. This approach provides more informative feedback to  
066    the agent, guiding its learning process. We treat the option price as a advanced exploration  
067    bonus for the agent, which encourages it to explore states randomly and discover potential  
068    rewards, even when immediate rewards are sparse.

069    However, we note that the BSM model assumes that the noise in the environment is Gaus-  
070    sian, which may not hold when the environment exhibits sudden changes or jumps. To ad-  
071    dress this limitation, we extend our approach by incorporating the Merton Jump-Diffusion  
072    Model, which accounts for jumps in the underlying asset price dynamics. This method how-  
073    ever, lead to the uncertainty in the convergence of the value estimates, thus we introduce an  
074    adapted decaying mechanism for the jump intensity parameter in the Merton model, which  
075    helps stabilize the learning process, guaranteeing the convergence of the value estimates.  
076

077    2 RELATED WORKS  
078

079    3 BSMRL  
080

081    3.1 BLACK-SCHOLES-MERTON MODEL  
082

083    In the field of financial mathematics, the Black-Scholes-Merton (BSM) model?? is a founda-  
084    tional framework for (European) option pricing. It assumes that the price of the underlying  
085    asset follows a geometric Brownian motion with constant volatility and drift. The BSM  
086    model provides a closed-form solution for European-style options.

087    They derived a partial differential equation (PDE) that the option price must satisfy, known  
088    as the Black-Scholes equation.

089    **Theorem 3.1** (Black-Scholes Equation). *The price of a European call option  $C(S, t)$  on a  
090    non-dividend-paying stock satisfies the following PDE:*

091    
$$\frac{\partial C}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 C}{\partial S^2} + rS \frac{\partial C}{\partial S} - rC = 0 \quad (1)$$
  
092

093    where  $C(S, t)$  is the price of the option at time  $t$  when the underlying asset price is  $S$ ,  $\sigma$  is  
094    the volatility of the underlying asset, and  $r$  is the risk-free interest rate.  
095

096    By applying Ito's Lemma and constructing a riskless portfolio, they eliminated the stochastic  
097    component and derived the pricing formula for European call options.  
098

099    **Theorem 3.2** (Black-Scholes-Merton Formula). *The price of a European call option  $C(S_t, t)$   
100    on a non-dividend-paying stock is given by:*

101    
$$C(S_t, t) = S_t \Phi(d_1) - K e^{-r(T-t)} \Phi(d_2) \quad (2)$$
  
102

103    where:  
104

$$d_1 = \frac{\ln(S_t/K) + (r + \frac{1}{2}\sigma^2)(T - t)}{\sigma\sqrt{T - t}}, \quad d_2 = d_1 - \sigma\sqrt{T - t} \quad (3)$$

105    Here,  $C(S_t, t)$  is the price of a European call option at time  $t$ ,  $S_t$  is the current price of  
106    the underlying asset,  $K$  is the strike price,  $r$  is the risk-free interest rate,  $\sigma$  is the volatility  
107    of the underlying asset,  $T$  is the time to maturity, and  $\Phi(\cdot)$  is the cumulative distribution  
108    function of the standard normal distribution.  
109

---

108 **Rationale** In Reinforcement Learning, the agent interacts with an environment to maximize cumulative rewards. However, in real-world scenarios, rewards can be sparse and delayed, making it challenging for agents to learn effective policies, leading to high variance in value estimates and slow convergence. To address this, we propose using the BSM model to shape rewards based on the agent’s current state and time-to-go, providing more informative feedback and guiding the agent’s learning process.

114 In a RL task, we have to estimate the value function  $V(s)$  or action-value function  $Q(s, a)$ ,  
115 usually within a certain period of time. For instance, in Monte-Carlo methods, we estimate  
116 the expected return over an episode, while in Temporal-Difference (TD) learning, we estimate  
117 the value function based on one-step transitions. This is analogous to option pricing,  
118 where the option’s value depends on the underlying asset price and time to maturity. These  
119 method often suffer from slow convergence rate due to sparse rewards, e.g., in website bug  
120 mining, the agent only receives a reward when a bug is found, which may take a long time.

121 By mimicing the idea of option pricing, we use the BSM formula and its derivatives to shape  
122 the expected accumulated reward, providing power for the agent to discover although the  
123 immediate reward is sparse. The method will provide a time-sensitive potential shaping  
124 reward, which can be viewed as an additional reward signal that guides the agent towards  
125 more promising states and actions.

126 However, such model assumes that the noise in the environment is Gaussian, which may not  
127 hold when the environment exhibits sudden changes or jumps. To address this limitation,  
128 we extend our approach by incorporating the Merton Jump-Diffusion Model, which accounts  
129 for jumps in the underlying asset price dynamics.

### 131 3.2 MERTON JUMP MODEL

133 In financial markets, asset prices often exhibit sudden and significant changes, known as  
134 jumps, which cannot be captured by the standard Black-Scholes model. To address this  
135 limitation, Robert C. Merton extended the Black-Scholes framework by incorporating jump  
136 processes into the asset price dynamics, leading to the Merton Jump-Diffusion Model. The  
137 Merton Jump-Diffusion Model assumes that the underlying asset price follows a stochastic  
138 process that combines both continuous diffusion and discrete jumps. The asset price  
139 dynamics under the Merton model can be described by the following stochastic differential  
140 equation (SDE):

$$dS_t = \mu S_t dt + \sigma S_t dW_t + J_t S_t dN_t \quad (4)$$

141 where:

- 143 •  $S_t$  is the asset price at time  $t$ .
- 144 •  $\mu$  is the drift rate of the asset price.
- 145 •  $\sigma$  is the volatility of the continuous component.
- 146 •  $W_t$  is a standard Brownian motion.
- 147 •  $N_t$  is a Poisson process with intensity  $\lambda$ , representing the number of jumps up to  
148 time  $t$ .
- 149 •  $J_t$  is the jump size, typically modeled as a log-normal random variable.

151 **Remark 3.1.** Such model can be also viewed as a Levy process, which generalizes Brownian  
152 motion by allowing for jumps.

154 The Merton Jump-Diffusion Model leads to a modified option pricing formula that accounts  
155 for the possibility of jumps in the underlying asset price.

156 **Theorem 3.3** (Merton Jump-Diffusion Option Pricing Formula). *The price of a European  
157 call option  $C(S_t, t)$  under the Merton Jump-Diffusion Model is given by:*

$$C(S_t, t) = \sum_{n=0}^{\infty} \frac{e^{-\lambda(T-t)} (\lambda(T-t))^n}{n!} C_{BS}(S_t, t; \sigma_n) \quad (5)$$

161 where:

- 
- 162     •  $C_{BS}(S_t, t; \sigma_n)$  is the Black-Scholes price of the option with adjusted volatility  $\sigma_n =$   
 163        $\sqrt{\sigma^2 + \frac{n\delta^2}{T-t}}$ , where  $\delta$  is the standard deviation of the jump size.  
 164  
 165     •  $\lambda$  is the jump intensity.  
 166  
 167     •  $T$  is the time to maturity.  
 168  
 169     •  $n$  is the number of jumps.  
 170  
 171

172  
 173     **Rationale** While the Black-Scholes model assumes continuous price movements, real-world environments often exhibit sudden changes or jumps, leading to fat-tailed reward distributions.  
 174  
 175

176     To better capture these dynamics, we propose using the Merton Jump-Diffusion Model for  
 177     reward shaping in Reinforcement Learning. By incorporating jump processes, the Merton  
 178     model provides a more accurate representation of environments with abrupt changes.  
 179

180  
 181     3.3 ADAPTED DECAYING FOR POISSON JUMP INTENSITY  
 182

183     We have assumed that the jump intensity  $\lambda$  is constant in the Merton model. However,  
 184     such assumption may lead to the overestimation of jump effects, resulting in non-convergent  
 185     value estimates. To address this, we introduce an adapted decaying mechanism for the jump  
 186     intensity  $\lambda$ .  
 187

188     We use the random counting process  $N(t)$  to model the number of jumps occurring up to  
 189     time  $t$ . When a significant jump is detected in the reward signal, we increase the jump  
 190     intensity  $\lambda$  to reflect the higher likelihood of future jumps. As the  $N(t) \sim P(\lambda)$ , then  
 191      $t \sim \Gamma(n, \lambda)$ , where  $n$  is the number of the observed jumps. After time  $T$ , where the option  
 192     is bound to expire, we may count how many jumps with intensity  $\lambda$ . If there are  
 193     too few jumps, then it means that the jump intensity is overestimated, thus we decay the  $\lambda$   
 194     accordingly. Specifically, we update  $\lambda$  as follows:  
 195

$$\lambda \leftarrow \lambda \cdot \exp(-\beta \cdot (N(T) - \lambda T)) \quad (6)$$

196     And we sample the observed time period  $T$  from  $\Gamma(1, \lambda)$ .  
 197  
 198

199     **Rationale** If in real environment, within times  $T$ , we observe less jumps than expected  
 200      $\lambda T$ , it indicates that the jump intensity  $\lambda$  is overestimated. To address this, we introduce  
 201     an adapted decaying mechanism for  $\lambda$ . By adjusting  $\lambda$  based on observed jump counts, we  
 202     ensure that the model remains responsive to actual environmental dynamics, thus stabilizing  
 203     the learning process and improving convergence of value estimates. Then, we sample  $T$  from  
 204      $\Gamma(1, \lambda)$  to reflect the updated jump intensity.  
 205  
 206

207     4 ALGORITHM FRAMEWORK  
 208  
 209

210     The overall algorithm framework for BSMRL using Merton Potential Shaping is outlined in  
 211     Algorithm 1. We introduce a method from ?, which allows us to decompose the reward into  
 212     a predictable component and a chaotic component, thus estimate the volatility rate of the  
 213     chaotic component, which will be used to compute the Merton potential shaping reward.  
 214  
 215

---

**Algorithm 1** BSMRL

```

216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269

1: Hyperparameters: Learning rate  $\alpha$ , Discount factor  $\gamma$ , Moving average rate  $\eta$  (for statistics), Maturity  $T$ , Risk-free rate  $r = -\ln \gamma$ , Jump mean  $\mu_J$ , Jump std  $\delta_J$ , Expansion terms  $K$ .
2: Initialize:
3:  $Q(s, a)$  arbitrarily.
4:  $\bar{R}(s, a) \leftarrow 0$  ▷ Predictable Reward Component
5:  $\Sigma_{chaos}^2(s, a) \leftarrow 0$  ▷ Chaotic Variance (Doob Volatility) [cite: 139]
6:  $\lambda_{jump}(s, a) \leftarrow \lambda_{init}$  ▷ Estimated Jump Intensity for Levy Process
7: Function BLACKSCHOLES( $S, K_{strike}, T, r, \sigma$ ):
8:  $d_1 \leftarrow \frac{\ln(S/K_{strike}) + (r + 0.5\sigma^2)T}{\sigma\sqrt{T}}$ 
9:  $d_2 \leftarrow d_1 - \sigma\sqrt{T}$ 
10: return  $S \cdot \Phi(d_1) - K_{strike}e^{-rT} \cdot \Phi(d_2)$ 
11: End Function
12: Function MERTONPRICE( $S_0, T, r, \sigma_{chaos}, \lambda, \mu_J, \delta_J$ ):
13:  $Price \leftarrow 0$ 
14:  $k \leftarrow e^{\mu_J + 0.5\delta_J^2} - 1$  ▷ Expected jump size
15:  $\lambda' \leftarrow \lambda(1 + k)$ 
16: for  $n = 0$  to  $K$  do
17:    $w_n \leftarrow \frac{e^{-\lambda'T}(\lambda'T)^n}{n!}$  ▷ Poisson weight for n jumps
18:    $\sigma_n \leftarrow \sqrt{\sigma_{chaos}^2 + \frac{n\delta_J^2}{T}}$  ▷ Effective volatility blending Doob noise & Jumps
19:    $r_n \leftarrow r - \lambda k + \frac{n \ln(1+k)}{T}$ 
20:    $Val_n \leftarrow \text{BLACKSCHOLES}(S_0, S_0, T, r_n, \sigma_n)$  ▷ ATM Call Option as Enhanced Value
21:    $Price \leftarrow Price + w_n \cdot Val_n$ 
22: end for
23: return  $Price$ 
24: End Function
25: loop
26: Initialize state  $s$ 
27: while  $s$  is not terminal do
28:   Choose action  $a$  (e.g.,  $\epsilon$ -greedy based on  $Q$ )
29:   Take action  $a$ , observe reward  $R_{obs}$  and next state  $s'$ 
30:    $\delta_{pred} \leftarrow R_{obs} - \bar{R}(s, a)$  ▷ Chaotic innovation
31:    $\bar{R}(s, a) \leftarrow \bar{R}(s, a) + \eta \cdot \delta_{pred}$  ▷ Update Predictable Component
32:    $\Sigma_{chaos}^2(s, a) \leftarrow (1 - \eta)\Sigma_{chaos}^2(s, a) + \eta \cdot (\delta_{pred})^2$  ▷ Update Chaotic Variance
33:    $\lambda_{jump}(s) \leftarrow \lambda_{jump}(s) \exp(-\beta(\frac{(V_t - V_{t-x})}{\lambda_{jump}(s)} - \lambda_{jump}(s)T))$ 
34:    $T \leftarrow \Gamma(1, \lambda_{jump}(s))$  ▷ Adapted Decaying for Jump Intensity
35:    $a'_{best} \leftarrow \arg \max_{a'} Q(s', a')$ 
36:    $S_{underlying} \leftarrow Q(s', a'_{best})$  ▷ Underlying asset is the naive value
37:    $\sigma_{input} \leftarrow \sqrt{\Sigma_{chaos}^2(s', a'_{best})}$  ▷ Use Doob Volatility as input
38:    $V_{enhanced}(s') \leftarrow \text{MERTONPRICE}(S_{underlying}, T, r, \sigma_{input}, \lambda_{jump}(s'), \mu_J, \delta_J)$ 
39:    $Y_{target} \leftarrow R_{obs} + \gamma \cdot V_{enhanced}(s')$  ▷ Maximize Option-Adjusted Return
40:    $Q(s, a) \leftarrow Q(s, a) + \alpha(Y_{target} - Q(s, a))$ 
41:    $s \leftarrow s'$ 
42: end while
43: end loop

```

---

---

270        5 THEORETICAL ANALYSIS  
271  
272        5.1 MDP FORMULATION FOR OPTIONRL  
273  
274        5.2 CONVERGENCE ANALYSIS  
275  
276        5.3 VARIANCE ANALYSIS  
277  
278        6 DISCUSSION AND FUTURE WORK  
279  
280        7 CONCLUSION  
281  
282        A PROOFS  
283  
284        B EXPERIMENT DETAILS  
285  
286  
287  
288  
289  
290  
291  
292  
293  
294  
295  
296  
297  
298  
299  
300  
301  
302  
303  
304  
305  
306  
307  
308  
309  
310  
311  
312  
313  
314  
315  
316  
317  
318  
319  
320  
321  
322  
323